

What is K8s?

Tuesday, June 21, 2022 9:59 AM

What is K8s?

Open Source container orchestration tool (dev by Google)

- Manages containers

- Manage containerized apps in different environments

What problems does K8s solve?

Trend from Monolith to Microservices increased usage of containers

- Managing 100s if not 1000s of containers

Orchestration Tools Features

- High Availability or no downtime

- Scalability

- Disaster recovery


Tuesday, June 21, 2022 10:05 AM

Node

Pod

Pods are ephemeral - can die easily


Ingress forwards to service, which sends to the pod



Built-in security mechanism is not enabled by default

You are responsible for backups/management of data

C.....'...k...DD...d...fNO...0k...STATELESS...'.d...NO...d...




Should be used to create stateful applications - Elasticsearch, MongoDB, MySQL

FOR STATEFULL APPLICATION

Ensures DB reads/writes synchronized

Deploying StatefulSet not easy

Common practice to host DB outside of K8s cluster & have STATELESS apps inside K8s cluster



K8s Architecture

Tuesday, June 21, 2022 10:31 AM

Node Processes

Worker Node

Worker machine in k8s cluster

Will have multiple Pods on it

Worker nodes do the actual work

3 processes must be installed on every Node:

Container Runtime - e.g. Docker

Kubelet - Process of K8s, interacts with both container & node

Responsible for applying configuration, starting the pod, & assigning resources

Kube Proxy - Forwards request from services to pods

Communication via Services - Load balancer that catches the request and forwards to respective pod

Has intelligent forwarding logic

Makes sure comms are performant, low overhead

Will forward to pod on same node before sending out to other nodes

Master Processes

How you interact with your K8s cluster

Managing processes are done by Master Nodes

4 process run on every master node

API Server - Cluster gateway, gets initial request / queries from cluster

Gatekeeper for authentication

Only 1 endpoint in to the cluster (good for security)

Scheduler - Schedules a new pod

Intelligent way of deciding which pod will get scheduled request

See your request & interpret number of resources you will need, then schedules on the node that has the resources to meet your Pod's needs

Scheduler just decides which Node a new pod should be scheduled

Kubelet receives request from the scheduler and actually executes the task

Controller Manager - Detects cluster state changes

Detects when pods die & tries to restore cluster state

Controller Manager -> Scheduler -> Kubelet

etcd - Key-value store of cluster state / cluster brain

Cluster changes get stored in key-value store

Retains information like:

What resources are available?

Did the cluster state change?

Is the cluster healthy?

Application data is NOT stored in etcd

K8s cluster typically has multiple Master nodes

Load balance API Server requests

Etcd forms distributed storage across master nodes

Example Cluster Set-up

Master Nodes need less resources than Worker Nodes since the worker nodes do the work

To **add new Master/Node server**:

1. Get new bare server

2. Install all the master/worker node processes
 3. Add it to the cluster
- Can infinitely scale horizontally

Minikube & Kubectl

Tuesday, June 21, 2022 10:52 AM

3 Ways to interact with K8s

UI
API
CLI

Minikube

Test / Local Cluster Setup

1 node K8s cluster, where all master / worker processes run on one node

Docker runtime pre-installed

Creates Virtual Box on your laptop & Node runs in that Virtual Box

For Testing purposes

Minikube has Kubectl as a dependency; will install both

Comes with "minikube" command line tool for basic interactions with cluster

Mainly for starting / stopping the cluster - management is done with Kubectl

Kubectl

Command Line Tool to Interact with your K8s cluster

Interacts directly with API Server installed on Master nodes

Most powerful method of interacting with K8s

Interacts with any K8s setup

Installation

Install Minikube (Mac, Linux and Windows): <https://bit.ly/38bLcJy>

Install Kubectl: <https://bit.ly/32bSI2Z>

Minikube Commands

Create & Start K8s & Select Hypervisor

`minikube start --vm-driver=hyperkit`

Also configures kubectl to use "minikube"

Get status of nodes in cluster

`kubectl get nodes`

`minikube status`

Get Version

`kubectl version`

Main Kubectl Commands

Tuesday, June 21, 2022 11:34 AM

List Kubernetes components

```
kubectl get nodes
kubectl get pod
kubectl get services
kubectl get deployment
kubectl get replicaset
```

Create a pod

```
kubectl create
    Pod is smallest unit, but you are creating deployment (abstraction over pods)
kubectl create deployment NAME --image=image
    Image is required
```

Edit a Deployment

```
kubectl edit deployment NAME
    Brings up auto-generated config file of deployment
    If config edited - will start a new pod, then remove the old one
    In practice, would be done within Kubernetes Configuration files
```

Get Application Logs

```
kubectl logs POD NAME
```

Get state changes

```
kubectl describe pod POD NAME
```

Create Bash Shell within the Container

```
kubectl exec -it POD NAME -- bin/bash
```

Delete a deployment

```
kubectl delete deployment NAME
```

Create deployment from Configuration File

```
kubectl apply -f CONFIG FILE
    Can also be used to apply changes that have been made to the configuration file
    K8s knows when to create or update deployment
```

Replicaset - manages replicas of the pod
Don't typically interact with



CRUD happens on Deployment level

CRUD commands

| | |
|-------------------|----------------------------------|
| Create deployment | kubectl create deployment [name] |
| Edit deployment | kubectl edit deployment [name] |
| Delete deployment | kubectl delete deployment [name] |

Status of different K8s components

kubectl get nodes | pod | services | replicaset | deployment

Debugging pods

| | |
|--------------------------|---|
| Log to console | kubectl logs (pod name) |
| Get interactive Terminal | kubectl exec -it (pod name) -- bin/bash |

Debugging pods

| | |
|--------------------------|---|
| Log to console | kubectl logs (pod name) |
| Get interactive Terminal | kubectl exec -it (pod name) -- bin/bash |

Use configuration file for CRUD

| | |
|--------------------------------|-------------------------------|
| Apply a configuration file | kubectl apply -f [file name] |
| Delete with configuration file | kubectl delete -f [file name] |

YAML Config File in K8s

Tuesday, June 21, 2022 11:54 AM

Every configuration has 3 parts

Metadata

Contains things like name / labels

Specification (spec)

Configuration to apply to deployment

Attributes specific to the kind of deployment

apiVersion / kind - base level

Declares what you want to create

Status

Automatically generated & added by Kubernetes

Desired state vs. Actual state

How you want it configured vs. how it's actually configured

Status info for actual comes from etcd

YAML Format

Configuration location

Configuration files stored with your code or it's own git repo for configs

Pod configuration

spec.template

Has it's own "metadata" and "spec" section

Blueprint for a pod

Connecting Components

Components are controlled with the "kind" variable

Can have a **kind** of Deployment or Service

Labels & Selectors

Establish the connection

Metadata contains the labels

Spec contains the selectors

Any key-value pair for a component

Deployment & Pods get their own labels

Connecting Services to Deployments

Within the service, leverage the "spec.selector" config to select deployments/pods that are part of the service

Ports in Service & Pod

Service will have a list of ports to listen externally and where to redirect internally

Some pods "containerPort" should line-up with the service's "targetPort"

Validation of Service / Deployment Connection

To check that a service is hooked into a specific endpoint

1. Get Metadata about service
 - a. `kubectl describe service SERVICE NAME`
2. Under "Endpoints", it should list the IP addresses / ports of the Pods it is pointing at
3. Get additional information about Pods
 - a. `kubectl get pod -o wide`

Validate Status

`kubectl get deployment DEPLOYMENT NAME -o yaml`



Demo Project MongoDB & MongoExpress

Wednesday, June 22, 2022 12:03 PM

Namespaces

Wednesday, June 22, 2022 11:24 AM

What is a Namespace?

Organize resources within Namespaces

Virtual cluster within a cluster

kubectl get namespaces

4 namespaces per Default

[kubernetes-dashboard](#)

Only with minikube

Will not be in standard cluster

[kube-system](#)

Do NOT create / modify in kube-system

System processes

Master / Kubectl process

[kube-public](#)

Publicly accessible data

ConfigMap which contains cluster information - accessible without authentication

kubectl cluster-info

[kube-node-lease](#)

Recent addition

Heartbeats of nodes

Each node has associated lease object in namespaces

Determines the availability of a node

[Default](#)

Resources you create are located here (unless specified otherwise)

Create a namespace

kubectl create namespace NAME

Can be created with a configuration file - better way

(kind of ConfigMap)

Why use Namespaces

[Default behavior- Everything is in one namespace](#)

Difficult to have an overview of components within a namespace - **Better to group resources in Namespaces**

Logically grouping resources within your cluster

Examples:

Database namespace for database resources

Monitoring namespace for monitoring resources

Elastic Stack namespace for Elastic stack (ELK)

According to Kubernetes documentation - Should not use for smaller projects or less than 10 users

[Conflicts: Many teams, same application](#)

Other teams can override deployment if they use the same name

If using Jenkins or automation, wouldn't know they overwrite a deployment

Each team can work in their own namespace without disrupting the other

[Resource Sharing: Staging & Development](#)

Deploy common resources once and reach from both environments

[Resource Sharing: Blue/Green Deployment](#)

Two versions of production in 1 cluster:

Current production (blue), next version of production (green)

Both current prod and next prod can leverage same resources

[Access & Resource limits on Namespaces](#)

Can give teams access to only their namespaces

Each team has its own, isolated environment

Can limit CPU, RAM, Storage per Namespace

Characteristics of Namespaces

[Can't access most resources from another namespace](#)

Each Namespace must define its own ConfigMap - even if it's own service

Can share services across namespaces

Can access service through "service_name.namespace" (e.g. mysql-service.database)

Some components can't be created within a namespace

Live globally within a cluster

Can't isolate them

Examples: Volume & Node

Can list components not in namespace: **kubectl api-resources --namespaced=false**

Create component in a Namespace

No Namespace defined - will be added to the "default" namespace

Apply at run-time: **kubectl apply -f mysql-configmap.yaml --namespace=my-namespace**

Or within the configuration file itself: "metadata.namespace"

In config file is better since it's documented & easier to apply in automation



Change active namespace (or default namespace)

Not natively possible with Kubernetes, but possible with tool called **kubens**

Kubectx & Kubens

<https://github.com/ahmetb/kubectx>

Tools to switch between contexts (clusters) and namespaces easily / faster

Ingress

Thursday, June 23, 2022 9:50 AM

- Git Repo: <https://bit.ly/3mJHVFc>
- Ingress Controllers: <https://bit.ly/32dfHe3>
- Ingress Controller Bare Metal: <https://bit.ly/3kYdmLB>

External Service vs. Ingress

External service would open external IP address and port

Ingress does **not** open external IP address and port

Instead leverages a domain

Ingress redirects traffic to internal service

External Service Configuration



Type is "LoadBalancer" & need to configure "nodePort"

Ingress Configuration

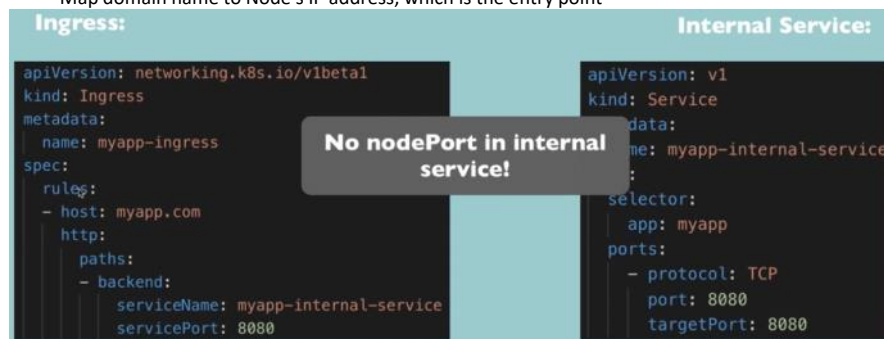


"http" key does not correlate to "https"

"host" needs to be:

Valid domain address

Map domain name to Node's IP address, which is the entry point



How to configure Ingress

Need an implementation for ingress, which is **ingress controller**

Ingress controller

Another Pod or set of pods that evaluate ingress rules

Evaluate all rules and manage redirections

Entrypoint to cluster

Many third-party implementations to choose from

e.g. K8s Nginx Ingress Controller

Cloud Environment

Cloud Load Balancer of choice would direct traffic to the ingress controller pod (one of many implementations)

Advantage of Cloud Load Balancer:

Don't have to implement load balancer yourself

Bare Metal Environment

Need to configure some kind of endpoint

Either inside or outside as a separate server

Example:

External Proxy Server

No server in K8s cluster is accessible from outside!

Endpoint forwards to ingress controller pod

Ingress Controller in Minikube

Installing in Minikube

minikube addons enable ingress

Automatically starts the K8s Nginx implementation of Ingress Controller

[See Nginx ingress controller running](#)

kubectl get pod -n kube-system

[See all components in Namespace](#)

kubectl get all -n kubernetes-dashboard

Configuring TLS Certificate

TLS Secret data needs to have tls.crt and tls.key key-value pairs

tls.crt & tls.key are file contents NOT paths/locations

Secret must be in the same namespace as the ingress component

The diagram illustrates the configuration of an Ingress resource to use a Secret for TLS certificates. On the left, the Ingress resource is defined with the following YAML:

```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: tls-example-ingress
spec:
  tls:
  - hosts:
    - myapp.com
    secretName: myapp-secret-tls
  rules:
  - host: myapp.com
    http:
      paths:
      - path: /
        backend:
          serviceName: myapp-internal-service
          servicePort: 8080
```

An arrow points from the `secretName: myapp-secret-tls` field in the Ingress spec to the Secret resource on the right. The Secret resource is defined as follows:

```
apiVersion: v1
kind: Secret
metadata:
  name: myapp-secret-tls
  namespace: default
data:
  tls.crt: base64 encoded cert
  tls.key: base64 encoded key
type: kubernetes.io/tls
```

Icons for 'ing' (Ingress) and 'secret' are shown next to their respective resource definitions.

Helm Package Manager

Thursday, June 23, 2022 11:04 AM

What is Helm?

Package manager for Kubernetes

Package YAML Files and Distribute them in public and private repositories

Helm Charts

Bundle of YAML files

Pushed to Helm Repository

helm install CHART NAME

Sharing Helm Charts

Searching from command line: `helm search <keyword>`

Or from Helm Hub

There are also private registries for Helm charts

Templating Engine

If deployment & service configs are almost the same:

You can define a common blueprint

Leverage dynamic values with placeholders

values.yaml file

`"{{ .Values.name }}"`

One yaml file for many microservices

Use cases

Practical for CI/CD - Build Yaml files & replace on the fly

Same applications across different environments

Deploy same setup in many environments

Helm Chart Structure

Directory Structure

mychart/ # Name of the chart

chart.yaml # Meta info about the chart (name version dependencies)

values.yaml # Values for the template files (default values)

charts/ # Chart dependencies (if it depends on other charts)

templates/ # Where template files are stored

...

`helm install <chartname>`

Template files will be filled with values from values.yaml

Values injection into template files

Default Values File: value.yaml

Define additional values files: `helm install --values=my-values.yaml <chartname>`

Any values not defined here, will pull from default value file

On the command line: `helm install --set key=value`

Release Management

Helm Version 2

Comes in two parts:

Client (helm CLI)

- helm install <chartname>

- Sends requests to Tiller

Server (Tiller)

- Where Kubernetes cluster is installed

- Keeps track of all chart executions

Changes are applied to existing deployment instead of creating a new one

Initially install chart: helm install <chartname>

Upgrade chart: helm upgrade <chartname>

Rollback changes: helm rollback <chartname>

[Downsides of Tiller](#)

Has too much power inside of K8s Cluster - Big Security Issue

Helm 3 removes Tiller - now a simple helm binary

Volumes / Persistence

Thursday, June 23, 2022 12:36 PM

Three tiers of Persistent Storage

Persistent Volume

physical or logical storage, location can be remote or local
Can also configure what type of access is required (accessMode)
Not associated with a namespace

Examples: AWS EBS, Local Hard disk

Persistent Volume Claim

claim made by developer stating how much storage is required for a pod
Associated with a namespace

Storage Classes

Provisions PV's dynamically when a PVC claims it
Abstracts underlying storage provider
Define the "**provisioner**" attribute for a specific storage provider for where to create PV's
Each storage backend has its own provisioner
Internal Provisioner - "kubernetes.io/" pre-fixed
Pre-defined provisioners native to Kubernetes
External Provisioner - provisioners not included in Kubernetes
Example: [NFS](#)

Define the "**parameters**" attribute for PV

Storage Class Usage

Requested with PVC that references storage class instead of persistent volume

1. Pod claims storage via PVC
2. PVC requests storage from StorageClass
3. SC creates PV that meets needs of the claim

[Kubernetes Documentation: <https://kubernetes.io/docs/concepts/storage/persistent-volumes/>]

ConfigMap & Secret

Local volumes, not created by PV or PVC, created by Kubernetes

Use cases

Need a file available for your pod
Need configuration file for your pod
Need certificate mounted for your pod

How it works

1. Create ConfigMap and/or Secret component
2. Mount that into your pod/container (same way as volume)

StatefulSet

Friday, June 24, 2022 11:05 AM

What is a StatefulSet?

K8s component used for Stateful applications

Stateful Applications - any applications that stores data to track state (e.g. databases)

Stateless Applications - don't keep record of state, each request is completely new

Best Practice to leverage Persistent storage along with Stateful applications

Deployments

Stateless Applications

Deployed using Deployment so that you can replacate your app as much as you want

Stateful Application

Deployed using StatefulSet components

Can replicate Pods, both manage pods based on an identical container specification

Configures storage in the same way as Stateless applications

Deployment Stateless vs. StatefulSet

Replicating stateful is more difficult due to more requirements

Stateless pods **can** be:

- Identical and interchangeable
- Created in random order with random hashes
- One Service that load balances to any Pod

Stateful pods **cannot** be:

- Created/deleted at same time
- Randomly addressed
- B/c replica Pods are not identical
 - Requires Pod Identity

Pod Identity with StatefulSet

Every Pod has it's own identifier - **Sticky identity** for each pod

Sticky identity allows each pod to retain state & role

Created from same specification, but not interchangeable

Persistent identifier across any re-scheduling

Why is identity necessary?

Multiple pods updating (writing) their own database would lead to **data inconsistency**

Master pod writes, worker/slave pods read

Each of the pods do not use the same physical storage

The each have a replica of the data that they can access themselves, leading to continuously syncing

Persistent Pod Identity will retain the Persistant storage, as well as Pod state through replacement

Important to use Remote Storage - needs to be accessible from all nodes!


Next Pod is only created if previous pod is up and running!

Delete StatefulSet or scaled down by N replicas, it will occur in the reverse order!

Each pod has it's own DNS Endpoint from a Service

Pod DNS name will be `${pod name}.${governing service domain}`

Governing service domain is the service name defined within the StatefulSet



When a Pod restarts: IP address changes, but name & endpoint stays the same

What you need to manage (not Kubernetes-managed)

Configuring the cloning & data synchronization

Make remote storage available

Managing and back-up

Stateful applications are not perfect for containerized environments

Kubernetes Services

Friday, June 24, 2022 11:58 AM

What is a Service?

Default Behavior without Service

Each pod gets an IP address from Node's range, but these are ephemeral b/c Pods are ephemeral (are destroyed frequently)

To see IP address of pod within the cluster: `kubectl get pod -o wide`

Service

Creates a stable IP address to access Pod(s)

Configures Loadbalancing

Loose coupling for comm within & outside the cluster

Spans all of the worker nodes

Attributes

Selectors tell the service which Pods to forward request to

"selector" attribute - key-value pairs for the labels of pods identify a set of pods

Service selector must match

targetPort attribute defines which port to send to on the Pod

When you create a service, K8s creates an **"Endpoint" object** that keeps track of which Pods are the endpoints of the service

Endpoint object will be the same name as the service

Get all endpoints: `kubectl get endpoints`

Service's **"port"** attribute is arbitrary, but **"targetPort"** attribute must match the port the container is listening at

Types of Services

ClusterIP Services

Default type of service

Service gets a static IP address & can receive requests on different ports

Multi-port Service

If you have multiple ports open within a service, you **must** "name" those ports

Head-less Services

Client wants to communicate with 1 specific Pod directly

Pods want to talk directly with a specific Pod

Not randomly selected

Use case: Stateful applications, like MySQL / MongoDB / Elasticsearch

Options for communicating with an individual Pod directly:

1. API Call to K8s API Server to get IP addresses (inefficient)
2. DNS Lookup
 - a. DNS Lookup for Service - returns single IP address (clusterIP)
 - b. Set **ClusterIP** to **"None"** - returns Pod IP address instead

Type Attributes

ClusterIP

Default, type not needed! **Internal** Service, only available from within the cluster

NodePort

External traffic has **access to a fixed port** on each worker node

Extension of the ClusterIP Service - Creates ClusterIP Service automatically

Defined in the "nodePort" attribute

Pre-defined range: 30000 - 32767

NodePort Services are not secure


NOT FOR PRODUCTION - Either configure Ingress or LoadBalancer for production environments

LoadBalancer

Service accessible **externally** through **cloud providers LoadBalancer**

Extension of the NodePort Service - NodePort & ClusterIP Service are created automatically

More secure since only one port is exposed to the LoadBalancer rather than outside



Can Combine Services if you need multiple options (communicate with cluster or directly with a single node)

